



Hybrid MPI & OpenMP Parallel Programming

MPI + OpenMP and other models on clusters of SMP nodes

Rolf Rabenseifner¹⁾

Rabenseifner@hlrs.de

Georg Hager²⁾

Georg.Hager@rrze.uni-erlangen.de

Gabriele Jost³⁾

gjost@tacc.utexas.edu

¹⁾ High Performance Computing Center (HLRS), University of Stuttgart, Germany

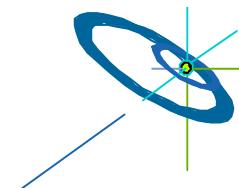
²⁾ Regional Computing Center (RRZE), University of Erlangen, Germany

³⁾ Texas Advanced Computing Center, The University of Texas at Austin, USA

HPC-Europa2 Virtual Surgery on “Hybrid MPI/OpenMP”

Dec. 3, 2010

(Summary from “Tutorial M02 at SC10, November 15, 2010, New Orleans, LA, USA”)



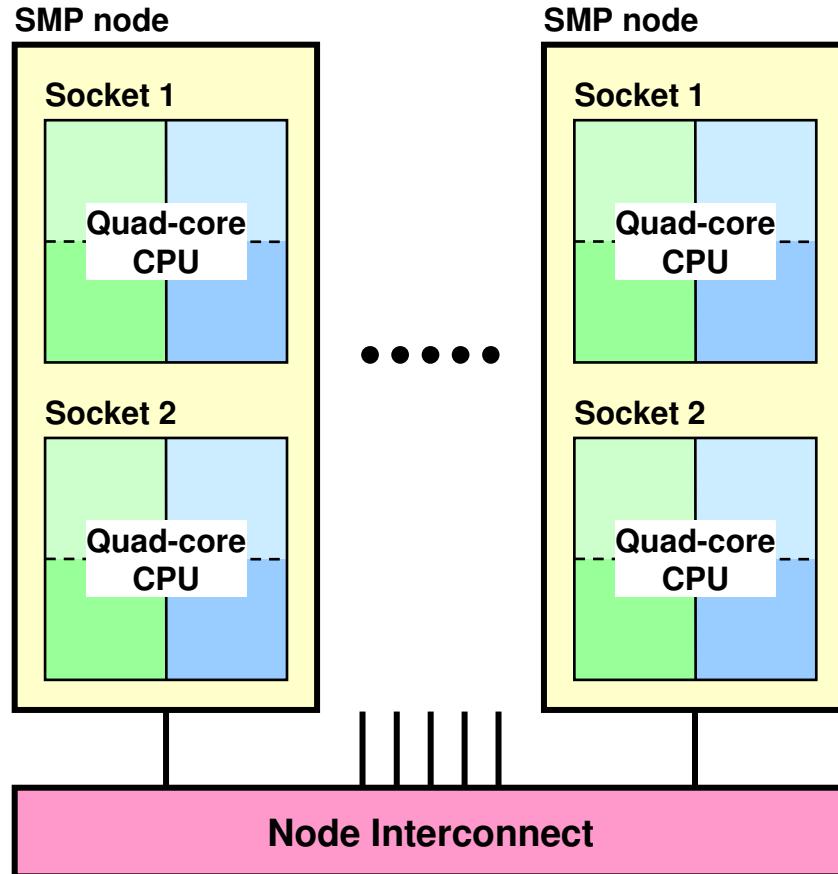
Hybrid Parallel Programming

Slide 1

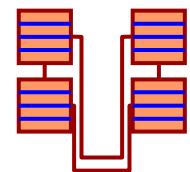
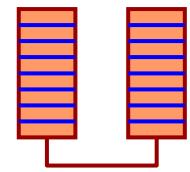
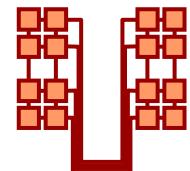
Höchstleistungsrechenzentrum Stuttgart



Motivation



- Which programming model is fastest?
- MPI everywhere?
- Fully hybrid MPI & OpenMP?
- Something between? (Mixed model)
- Often hybrid programming **slower** than pure MPI
 - Examples, Reasons, ...



?

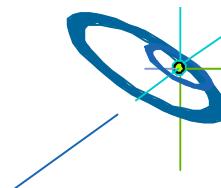
Outline

- **Programming models on clusters of SMP nodes**
 - Case Studies / pure MPI vs hybrid MPI+OpenMP
 - Mismatch Problems
 - Opportunities:
Application categories that can benefit from hybrid parallelization
 - Conclusion

Slides are available from

www.hlrs.de/people/rabenseifner

→ List of publications → International teaching



Parallel Programming Models on Hybrid Platforms

pure MPI
one MPI process
on each core

hybrid MPI+OpenMP
MPI: inter-node communication
OpenMP: inside of each SMP node

OpenMP only
distributed virtual
shared memory

No overlap of Comm. + Comp.
MPI only outside of parallel regions
of the numerical application code

Overlapping Comm. + Comp.
MPI communication by one or a few threads
while other threads are computing

Masteronly
MPI only outside
of parallel regions



Pure MPI

pure MPI
one MPI process
on each core

Advantages

- No modifications on existing MPI codes
- MPI library need not to support multiple threads

Major problems

- Does MPI library uses internally different protocols?
 - Shared memory inside of the SMP nodes
 - Network communication between the nodes
- Does application topology fit on hardware topology?
- Unnecessary MPI-communication inside of SMP nodes!

Discussed
in detail later on
in the section
**Mismatch
Problems**



Hybrid Masteronly

Masteronly
 MPI only outside
 of parallel regions

Advantages

- No message passing inside of the SMP nodes
- No topology problem

```

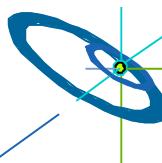
for (iteration ....)
{
  #pragma omp parallel
    numerical code
  /*end omp parallel */

  /* on master thread only */
  MPI_Send (original data
    to halo areas
    in other SMP nodes)
  MPI_Recv (halo data
    from the neighbors)
} /*end for loop

```

Major Problems

- All other threads are sleeping while master thread communicates!
- Which inter-node bandwidth?
- MPI-lib must support at least MPI_THREAD_FUNNELED

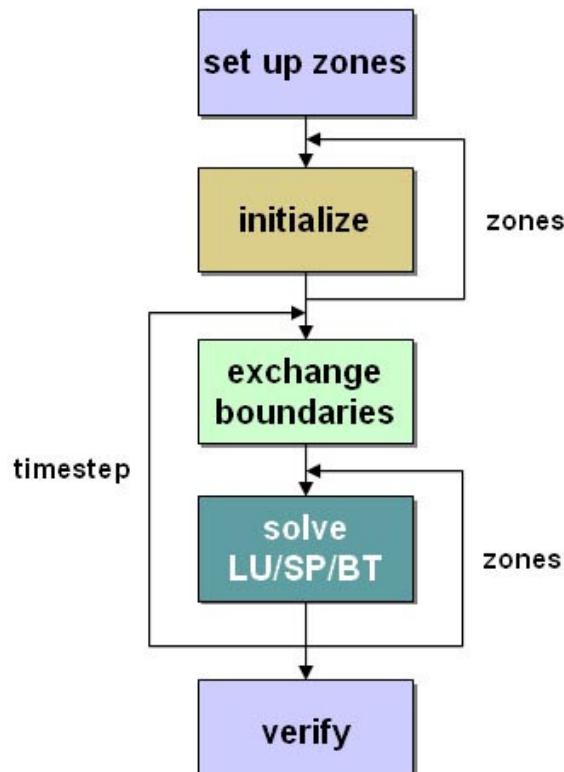


Outline

- Programming models on clusters of SMP nodes
- **Case Studies / pure MPI vs hybrid MPI+OpenMP**
- Mismatch Problems
- Opportunities:
Application categories that can benefit from hybrid parallelization
- Conclusion



The Multi-Zone NAS Parallel Benchmarks

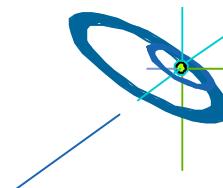


	MPI/OpenMP	MLP	Nested OpenMP
Time step	sequential	sequential	sequential
inter-zones	MPI Processes	MLP Processes	OpenMP
exchange boundaries	Call MPI	data copy+ sync.	OpenMP
intra-zones	OpenMP	OpenMP	OpenMP

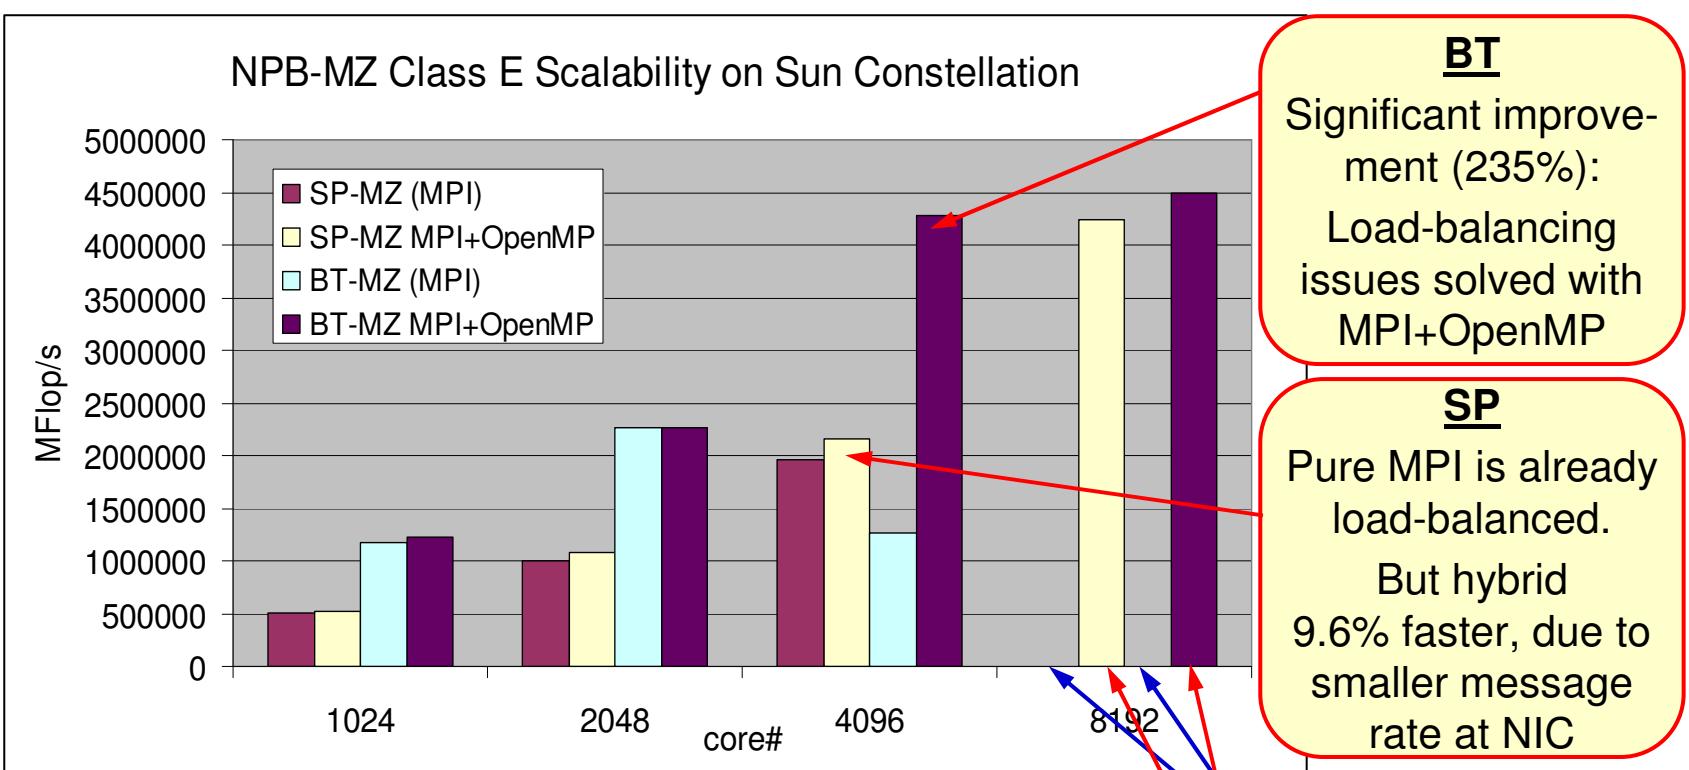
- Multi-zone versions of the NAS Parallel Benchmarks LU, SP, and BT
- Two hybrid sample implementations
- Load balance heuristics part of sample codes
- www.nas.nasa.gov/Resources/Software/software.html

The Multi-Zone NAS Parallel Benchmarks

- Aggregate sizes:
 - Class D: 1632 x 1216 x 34 grid points
 - Class E: 4224 x 3456 x 92 grid points
 - **BT-MZ: (Block tridiagonal simulated CFD application)**
 - Alternative Directions Implicit (ADI) method
 - #Zones: 1024 (D), 4096 (E)
 - Size of the zones varies widely:
 - large/small about 20
 - requires multi-level parallelism to achieve a good load-balance
 - **SP-MZ: (Scalar Pentadiagonal simulated CFD application)**
 - #Zones: 1024 (D), 4096 (E)
 - Size of zones identical
 - no load-balancing required
- Expectations:
- Pure MPI: Load-balancing problems!
 Good candidate for MPI+OpenMP
- Load-balanced on MPI level: Pure MPI should perform best

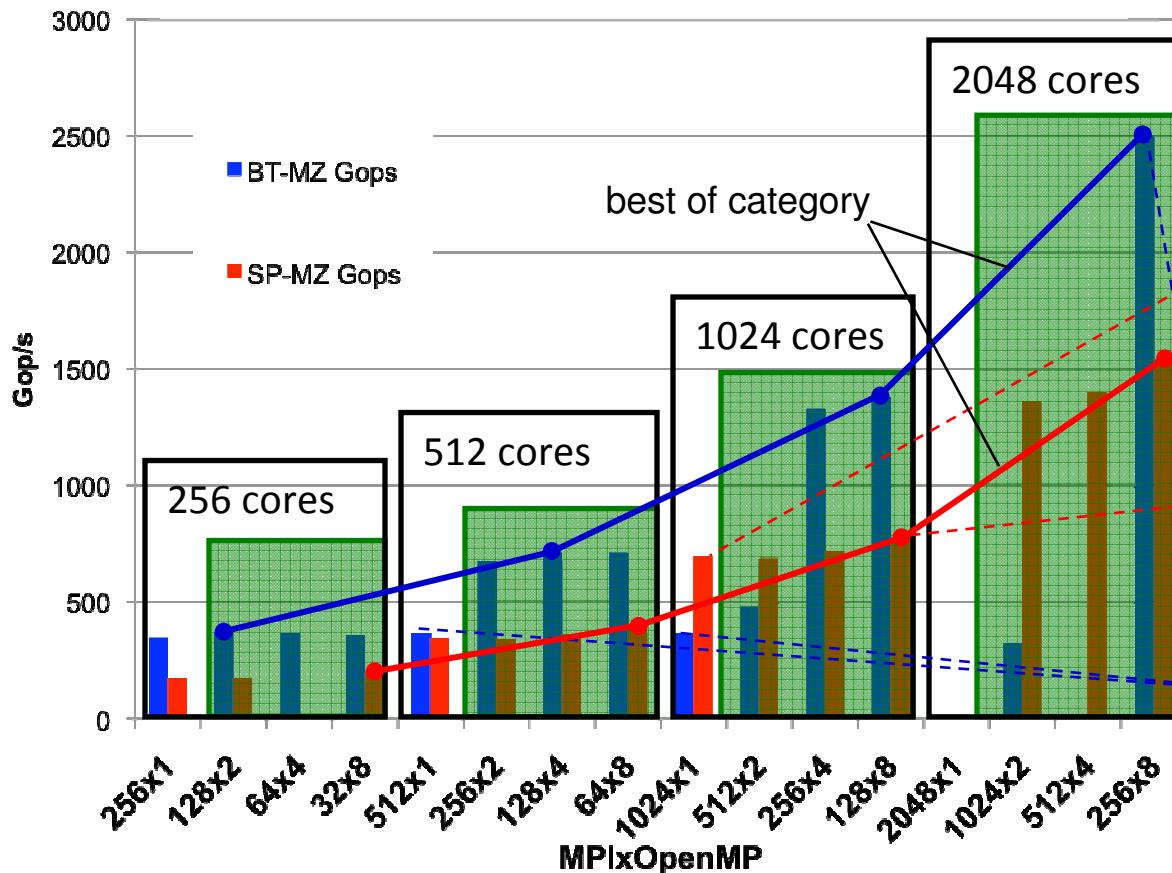


SUN: NPB-MZ Class E Scalability on Ranger



- Scalability in Mflops
- MPI/OpenMP outperforms pure MPI
- Use of numactl essential to achieve scalability

Cray XT5: NPB-MZ Class D Scalability



Results reported for
Class D on 256-2048 cores

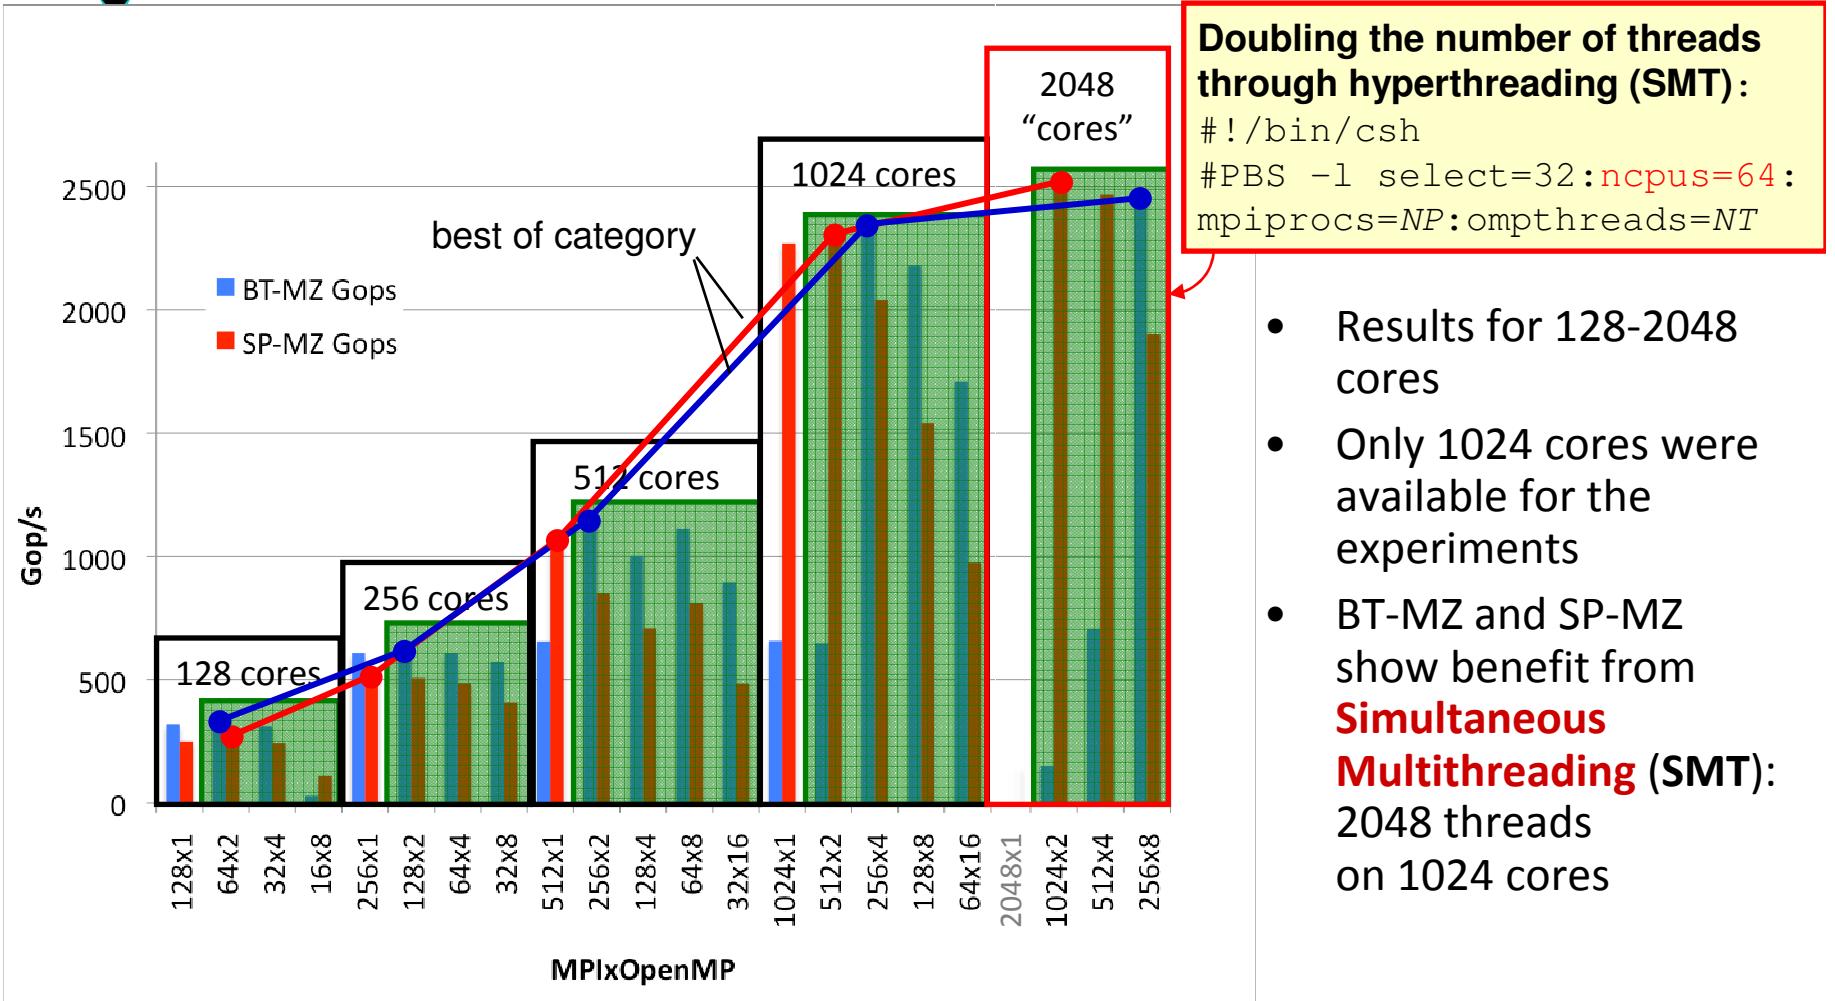
Expected:
#MPI processes limited

- SP-MZ pure MPI scales up to 1024 cores
- SP-MZ MPI/OpenMP scales to 2048 cores
- SP-MZ MPI/OpenMP outperforms pure MPI for 1024 cores
- BT-MZ MPI does not scale
- BT-MZ MPI/OpenMP scales to 2048 cores, outperforms pure MPI

Unexpected!

Expected: Load-
Imbalance for
pure MPI

NPB-MZ Class D on IBM Power 6: Exploiting SMT for 2048 Core Results



Intra-node MPI characteristics: IMB Ping-Pong benchmark

- Code (to be run on 2 processors):

```

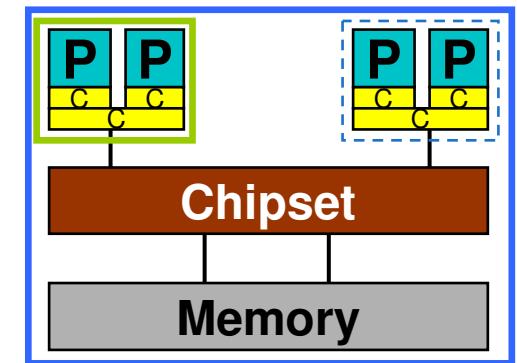
wc = MPI_WTIME()

do i=1,NREPEAT

  if(rank.eq.0) then
    MPI_SEND(buffer,N,MPI_BYTE,1,0,MPI_COMM_WORLD,ierr)
    MPI_RECV(buffer,N,MPI_BYTE,1,0,MPI_COMM_WORLD, &
              status,ierr)
  else
    MPI_RECV(...)
    MPI_SEND(...)
  endif

enddo

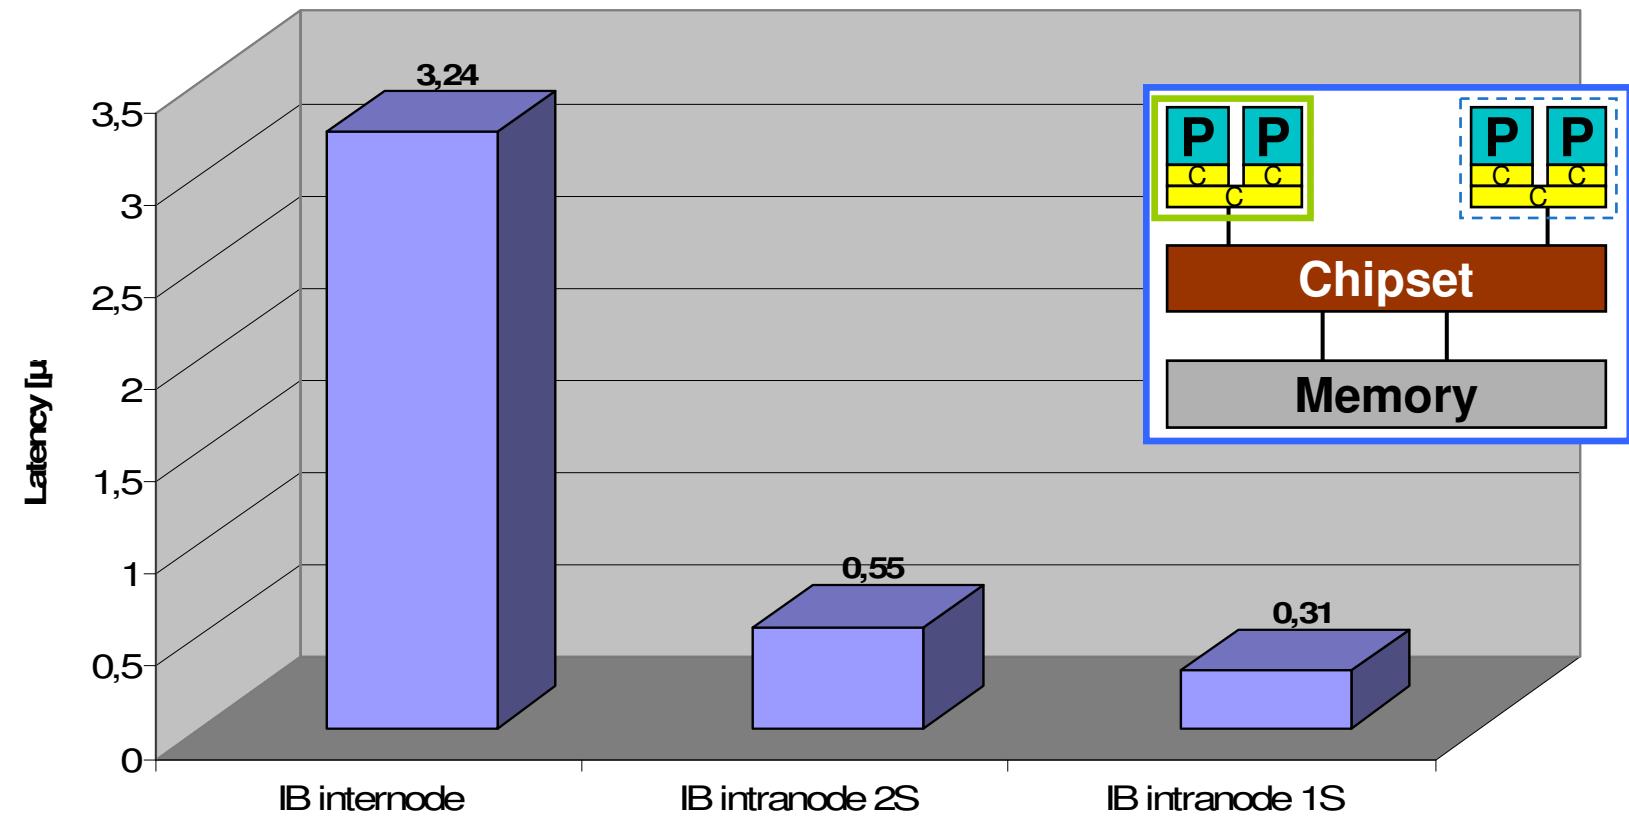
wc = MPI_WTIME() - wc
  
```



- Intranode (1S): mpirun -np 2 -pin "1 3" ./a.out
- Intranode (2S): mpirun -np 2 -pin "2 3" ./a.out
- Internode: mpirun -np 2 -pernode ./a.out

IMB Ping-Pong: Latency

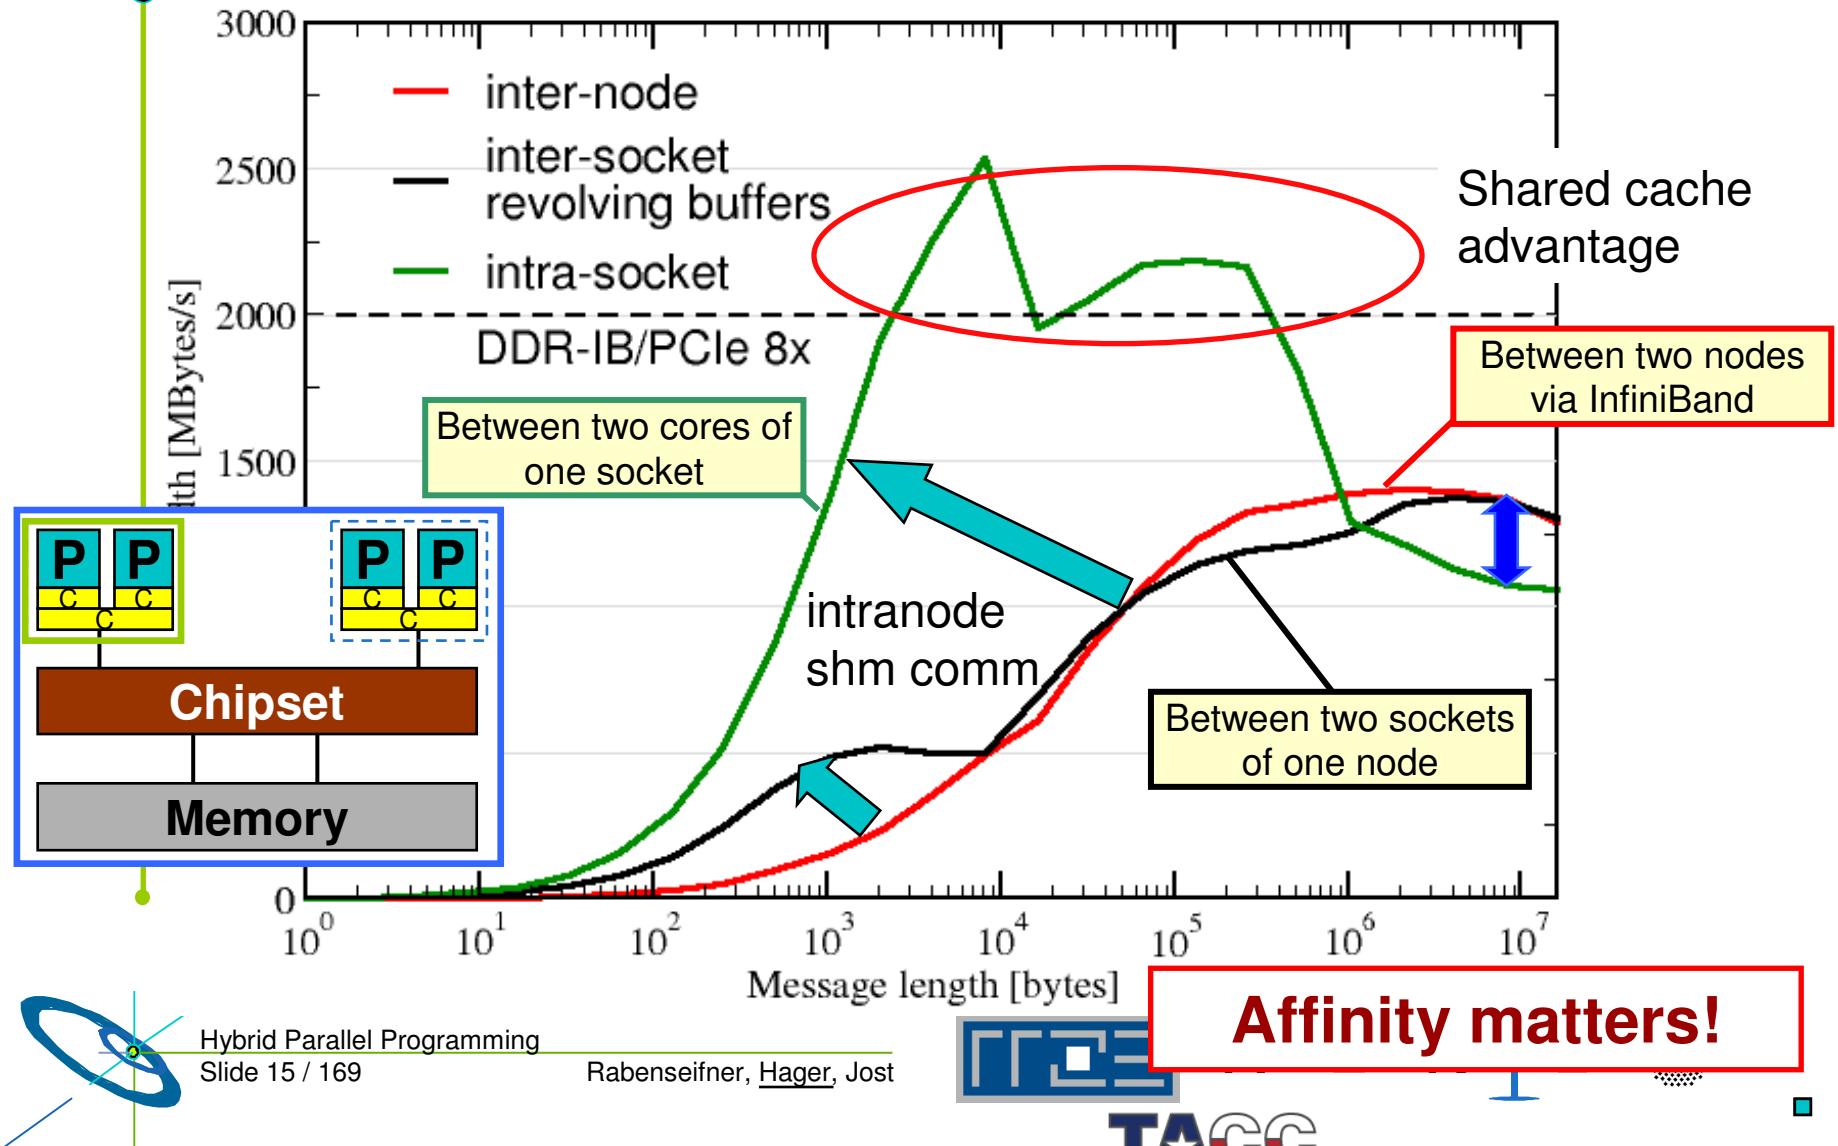
Intra-node vs. Inter-node on Woodcrest DDR-IB cluster (Intel MPI 3.1)



Affinity matters!

IMB Ping-Pong: Bandwidth Characteristics

Intra-node vs. Inter-node on Woodcrest DDR-IB cluster (Intel MPI 3.1)



Thread/Process Affinity (“Pinning”)

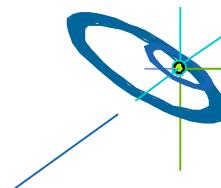
- Highly OS-dependent system calls
 - But available on all systems
 - Linux: `sched_setaffinity()`, PLPA (see below) → `hwloc`
 - Solaris: `processor_bind()`
 - Windows: `SetThreadAffinityMask()`
 - ...
- Support for “semi-automatic” pinning in some compilers/environments
 - Intel compilers > V9.1 (`KMP_AFFINITY` environment variable)
 - Pathscale
 - SGI Altix `dplace` (works with logical CPU numbers!)
 - Generic Linux: `taskset`, `numactl`, `likwid-pin`
- Affinity awareness in MPI libraries
 - SGI MPT
 - OpenMPI
 - Intel MPI
 - ...

Used on SUN Ranger slides



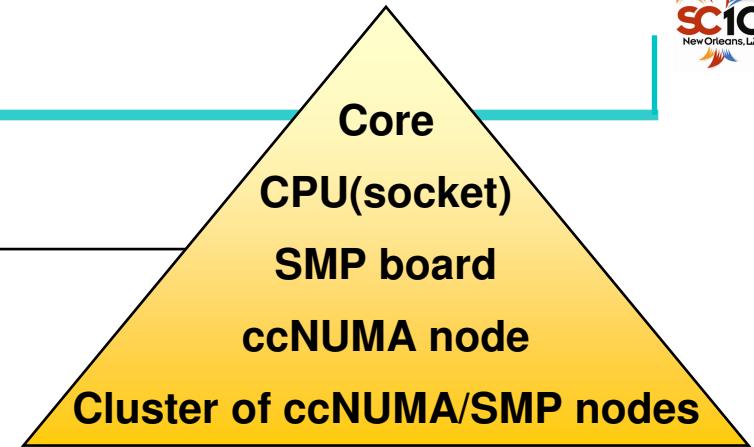
Outline

- Programming models on clusters of SMP nodes
- Case Studies / pure MPI vs hybrid MPI+OpenMP
- **Mismatch Problems**
- Opportunities:
Application categories that can benefit from hybrid parallelization
- Conclusion



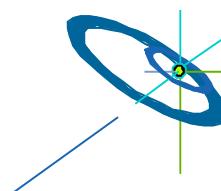
Mismatch Problems

- None of the programming models fits to the hierarchical hardware (cluster of SMP nodes)
- Several mismatch problems → following slides
- Benefit through hybrid programming → Opportunities, see next section
- Quantitative implications → depends on your application



Examples:	No.1	No.2
Benefit through hybrid (see next section)	30%	10%
<u>Loss by mismatch problems</u>	-10%	-25%
Total	+20%	-15%

In most cases:
Both categories!



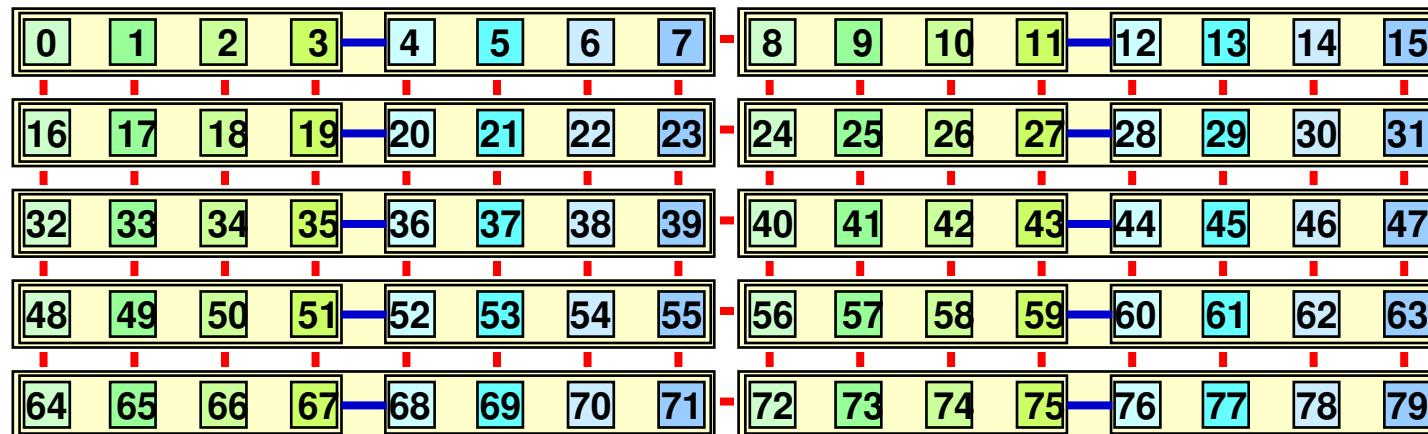
The Topology Problem with

pure MPI

one MPI process
on each core

Application example on 80 cores:

- Cartesian application with $5 \times 16 = 80$ sub-domains
- On system with $10 \times$ dual socket \times quad-core



- + 17 x inter-node connections per node
- 1 x inter-socket connection per node

Sequential ranking of
`MPI_COMM_WORLD`

Does it matter?

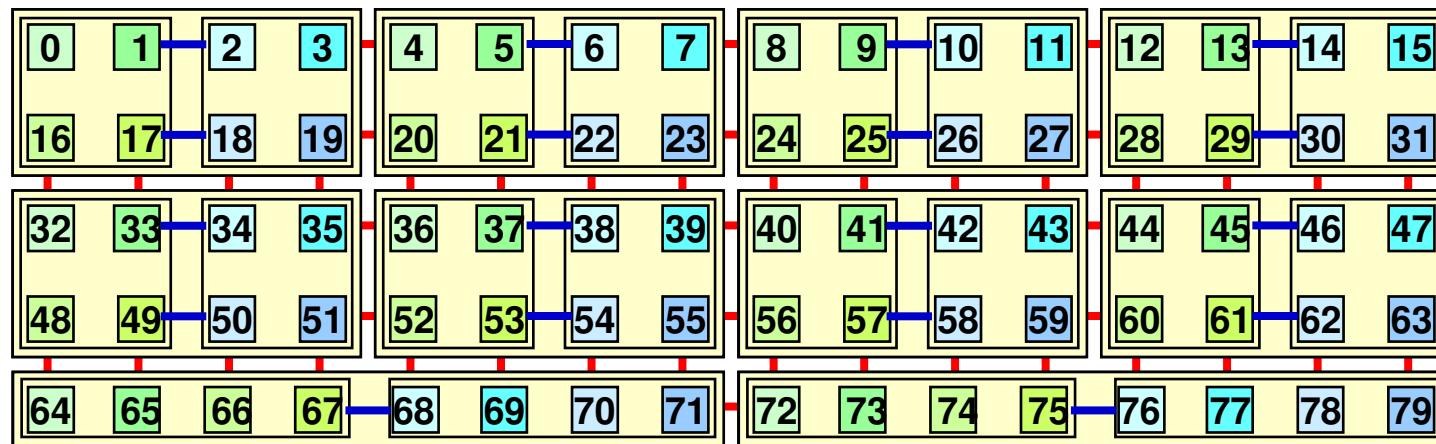
The Topology Problem with

pure MPI

one MPI process
on each core

Application example on 80 cores:

- Cartesian application with $5 \times 16 = 80$ sub-domains
- On system with $10 \times$ dual socket \times quad-core



- + 12 x inter-node connections per node
- + 2 x inter-socket connection per node

Two levels of
domain decomposition

Good affinity of cores to thread ranks



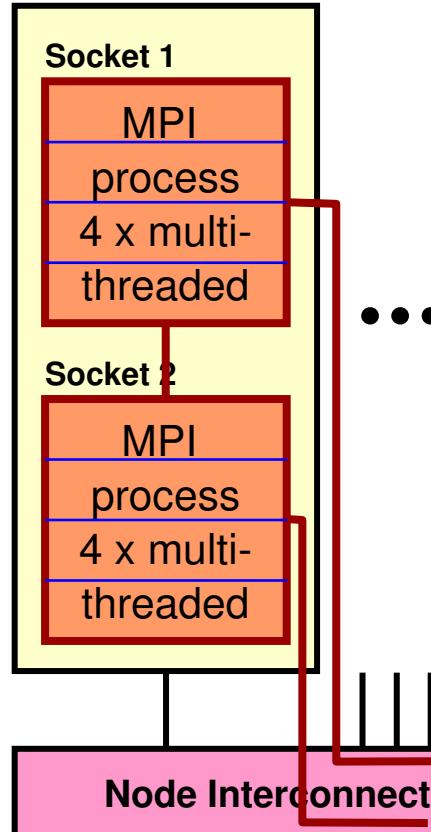
The Mapping Problem with mixed model

pure MPI
&

hybrid MPI+OpenMP

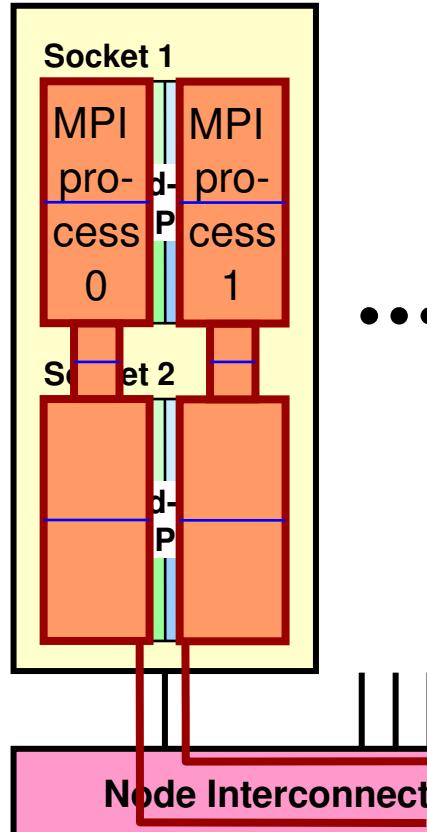
Do we have this?

SMP node



... or that?

SMP node



Several multi-threaded MPI process per SMP node:

Problem

- Where are your processes and threads really located?

Solutions:

- Depends on your platform,
- e.g., with **numactl**

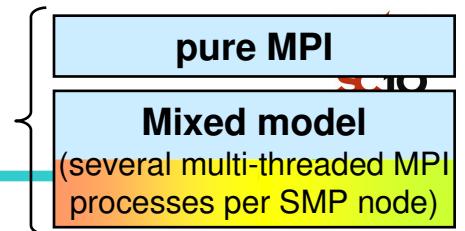
→ Case study on
Sun Constellation Cluster
Ranger
with BT-MZ and SP-MZ

Further questions:

- Where is the NIC¹⁾ located?
- Which cores share caches?



Unnecessary intra-node communication



Problem:

- If several MPI process on each SMP node
→ unnecessary intra-node communication

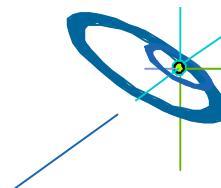
Solution:

- Only one MPI process per SMP node

Remarks:

- MPI library must use appropriate fabrics / protocol for intra-node communication
- Intra-node bandwidth higher than inter-node bandwidth
→ problem may be small
- MPI implementation may cause unnecessary data copying
→ waste of memory bandwidth

Quality aspects
of the MPI library



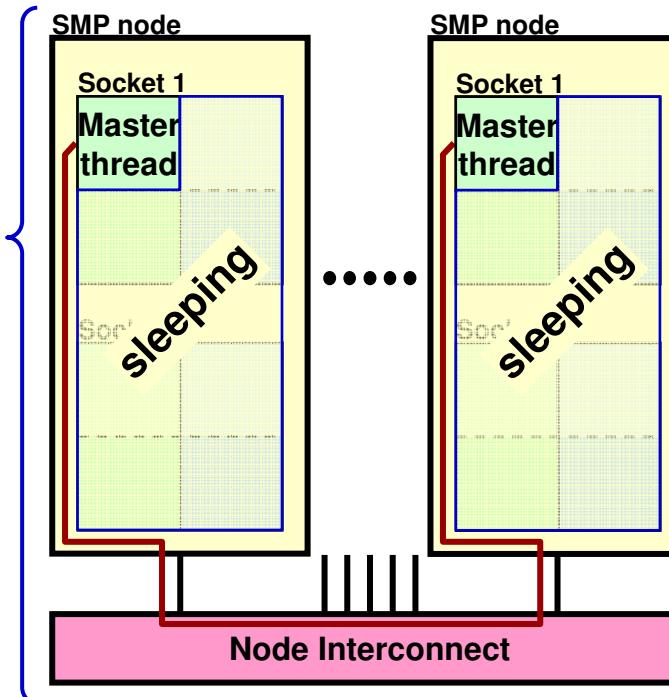
Sleeping threads and network saturation

with **Masteronly**

MPI only outside of parallel regions

```
for (iteration ....)
{
  #pragma omp parallel
  numerical code
  /*end omp parallel */

  /* on master thread only */
  MPI_Send (original data
            to halo areas
            in other SMP nodes)
  MPI_Recv (halo data
            from the neighbors)
} /*end for loop
```



Problem 1:

- Can the master thread saturate the network?

Solution:

- If not, use mixed model
- i.e., several MPI processes per SMP node

Problem 2:

- Sleeping threads are wasting CPU time

Solution:

- Overlapping of computation and communication

Problem 1&2 together:

- Producing more idle time through lousy bandwidth of master thread



OpenMP: Additional Overhead & Pitfalls

- Using OpenMP
 - may prohibit compiler optimization
 - **may cause significant loss of computational performance**
- Thread fork / join overhead
- On ccNUMA SMP nodes:
 - **Loss of performance due to missing memory page locality or missing first touch strategy**
 - E.g. with the masteronly scheme:
 - One thread produces data
 - Master thread sends the data with MPI
 - data may be internally communicated from one memory to the other one
- Amdahl's law for each level of parallelism
- Using MPI-parallel application libraries? → Are they prepared for hybrid?



Overlapping Communication and Computation

MPI communication by one or a few threads while other threads are computing

Three problems:

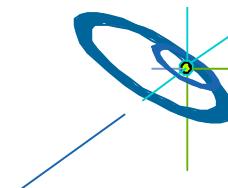
- the application problem:
 - one must separate application into:
 - code that can run before the halo data is received
 - code that needs halo data

→ very hard to do !!!
- the thread-rank problem:
 - comm. / comp. via thread-rank
 - cannot use work-sharing directives

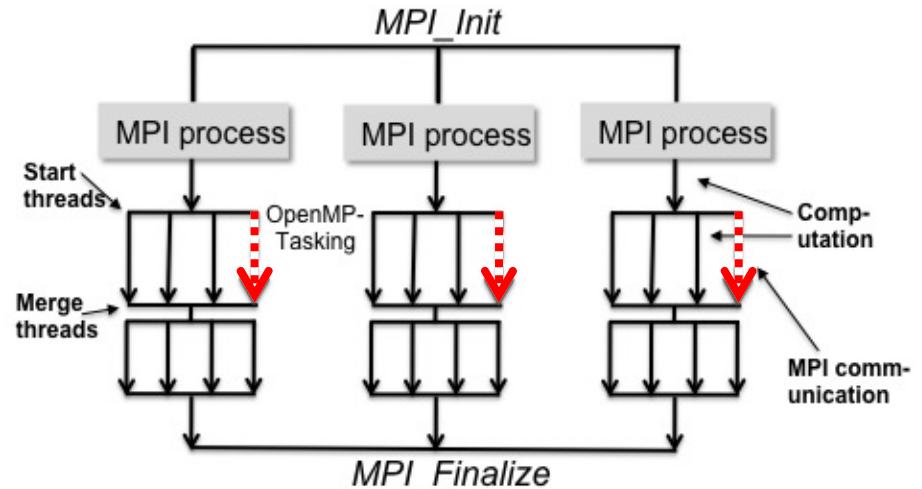
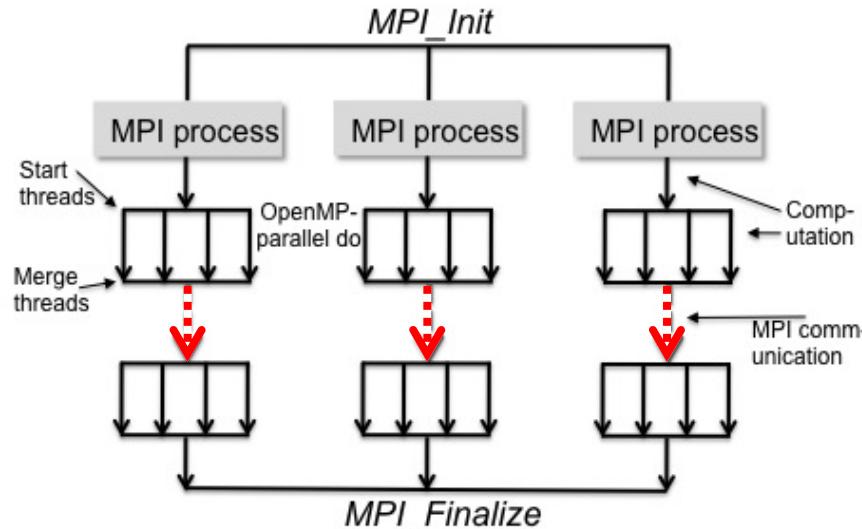
→ loss of major OpenMP support
 (see next slide)
- the load balancing problem

```

if (my_thread_rank < 1) {
  MPI_Send/Recv.....
} else {
  my_range = (high-low-1) / (num_threads-1) + 1;
  my_low = low + (my_thread_rank+1)*my_range;
  my_high=high+ (my_thread_rank+1+1)*my_range;
  my_high = max(high, my_high)
  for (i=my_low; i<my_high; i++) {
    ....
  }
}
  
```



Overlapping: Using OpenMP tasks



NEW OpenMP Tasking Model gives a new way to achieve more parallelism from hybrid computation.

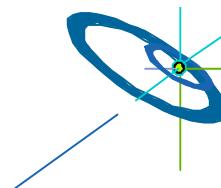
Alice Koniges et al.:
Application Acceleration on Current and Future Cray Platforms.
Proceedings, CUG 2010, Edinburgh, GB, May 24-27, 2010.

Courtesy of Alice Koniges, NERSC, LBNL



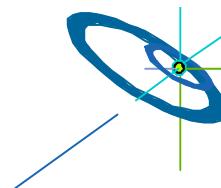
No silver bullet

- The analyzed programming models do **not** fit on hybrid architectures
 - whether drawbacks are minor or major
 - **depends on applications' needs**
 - But there are major opportunities → next section
 - In the NPB-MZ case-studies
 - We tried to use optimal parallel environment
 - **for pure MPI**
 - **for hybrid MPI+OpenMP**
 - i.e., the developers of the MZ codes and we tried to minimize the mismatch problems
- the opportunities in next section dominated the comparisons



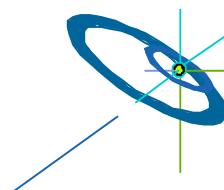
Outline

- Programming models on clusters of SMP nodes
- Case Studies / pure MPI vs hybrid MPI+OpenMP
- Mismatch Problems
- **Opportunities:**
Application categories that can benefit from hybrid parallelization
- Conclusion



Nested Parallelism

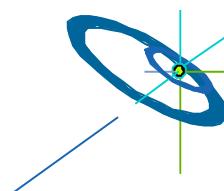
- Example NPB: BT-MZ (Block tridiagonal simulated CFD application)
 - Outer loop:
 - limited number of zones → limited parallelism
 - zones with different workload → speedup < $\frac{\text{Sum of workload of all zones}}{\text{Max workload of a zone}}$
 - Inner loop:
 - OpenMP parallelized (static schedule)
 - Not suitable for distributed memory parallelization
- Principles:
 - Limited parallelism on outer level
 - Additional inner level of parallelism
 - Inner level not suitable for MPI
 - Inner level may be suitable for static OpenMP worksharing



Load-Balancing (on same or different level of parallelism)

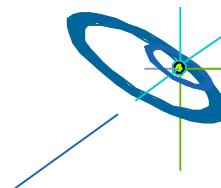


- OpenMP enables
 - Cheap **dynamic** and **guided** load-balancing
 - Just a parallelization option (clause on omp for / do directive)
 - Without additional software effort
 - Without explicit data movement
 - On MPI level
 - **Dynamic load balancing** requires moving of parts of the data structure through the network
 - Significant runtime overhead
 - Complicated software / therefore not implemented
 - **MPI & OpenMP**
 - Simple static load-balancing on MPI level, dynamic or guided on OpenMP level

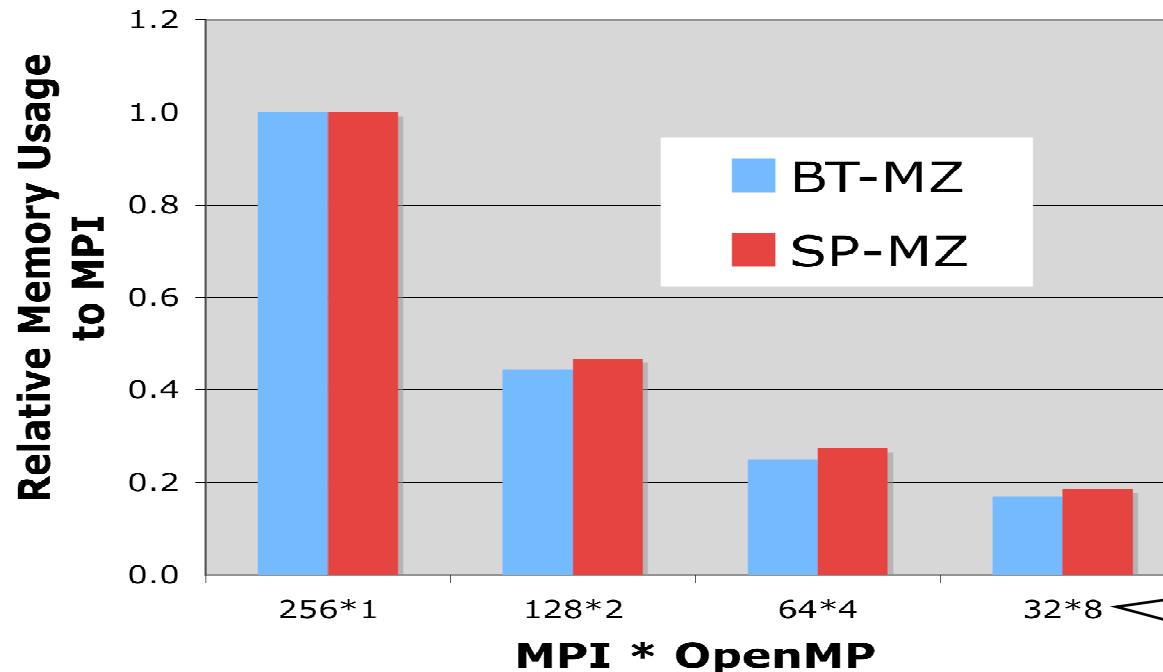


Memory consumption

- Shared nothing
 - Heroic theory
 - In practice: Some data is duplicated
- **MPI & OpenMP**
With n threads per MPI process:
 - Duplicated data may be reduced by factor n



Case study: MPI+OpenMP memory usage of NPB



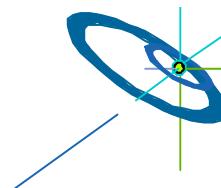
Using more OpenMP threads could reduce the memory usage **substantially**, up to **five** times on Hopper Cray XT5 (eight-core nodes).

Always same number of cores

Hongzhang Shan, Haoqiang Jin, Karl Fuerlinger,
 Alice Koniges, Nicholas J. Wright:
Analyzing the Effect of Different Programming Models Upon Performance and Memory Usage on Cray XT5 Platforms.
 Proceedings, CUG 2010, Edinburgh, GB, May 24-27, 2010.

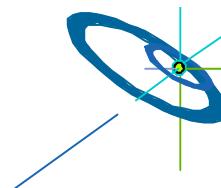
How many threads per MPI process?

- SMP node = with **m sockets** and **n cores/socket**
- How many threads (i.e., cores) per MPI process?
 - Too many threads per MPI process
 - overlapping of MPI and computation may be necessary,
 - some NICs unused?
 - Too few threads
 - too much memory consumption (see previous slides)
- Optimum
 - somewhere between 1 and $m \times n$ threads per MPI process,
 - Typically:
 - Optimum = n , i.e., 1 MPI process per socket
 - Sometimes = $n/2$ i.e., 2 MPI processes per socket
 - Seldom = $2n$, i.e., each MPI process on 2 sockets



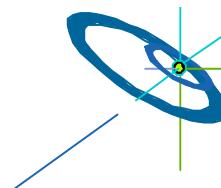
To overcome MPI scaling problems

- Reduced number of MPI messages, }
reduced aggregated message size } compared to pure MPI
- MPI has a few scaling problems
 - Handling of more than 10,000 MPI processes
 - Irregular Collectives: MPI_....v(), e.g. MPI_Gatherv()
 - Scaling applications should not use MPI_....v() routines
 - MPI-2.1 Graph topology (MPI_Graph_create)
 - MPI-2.2 MPI_Dist_graph_create_adjacent
 - Creation of sub-communicators with MPI_Comm_create
 - MPI-2.2 introduces a new scaling meaning of MPI_Comm_create
 - ... see P. Balaji, et al.: **MPI on a Million Processors**. Proceedings EuroPVM/MPI 2009.
- Hybrid programming reduces all these problems (due to a smaller number of processes)



Summary: Opportunities of hybrid parallelization (MPI & OpenMP)

- Nested Parallelism
 - Outer loop with MPI / inner loop with OpenMP
- Load-Balancing
 - Using OpenMP ***dynamic*** and ***guided*** worksharing
- Memory consumption
 - Significantly reduction of replicated data on MPI level
- Opportunities, if MPI speedup is limited due to algorithmic problem
 - Significantly reduced number of MPI processes
- Reduced MPI scaling problems
 - Significantly reduced number of MPI processes



Conclusions

- Future hardware will be more complicated
 - Heterogeneous → GPU, FPGA, ...
 - ccNUMA quality may be lost on cluster nodes
 -
- High-end programming → more complex
- Medium number of cores → programming will be more simple (provided that **#cores / SMP-node** will not shrink)
- **MPI+OpenMP → work horse on large systems**
- Pure MPI → still on smaller cluster
- OpenMP → on large ccNUMA nodes
(but not with shared virtual memory systems (e.g., ClusterOpenMP))



Thank you for your interest

→ Next talk by Rainer Keller

